

610 HW7

Joe Stoica and Corey Maxedon

11/10/2019

1

No output required.

2

```
get_inv_cov <- function(lambda, data=df) {  
  p = ncol(data)  
  theta = Variable(rows = p, cols = p)  
  S <- cov(data)  
  objective = Minimize(-log_det(theta) + matrix_trace(S %*% theta) + lambda * sum(abs(theta)))  
  problem = Problem(objective)  
  result = psolve(problem)  
  return(result$getValue(theta))  
}  
  
get_inv_cov(5)
```

```
##           [,1]           [,2]           [,3]           [,4]           [,5]  
## [1,] 1.413040e-01 5.678242e-08 5.918907e-08 5.672798e-08 5.780301e-08  
## [2,] 2.507180e-07 1.823866e-01 -5.151694e-08 -5.192415e-08 -4.773454e-08  
## [3,] 2.646992e-07 -1.719967e-07 1.808716e-01 -5.529544e-08 -5.006357e-08  
## [4,] 2.597271e-07 -1.727238e-07 -1.834903e-07 1.808871e-01 -4.830552e-08  
## [5,] 2.423702e-07 -1.548657e-07 -1.629256e-07 -1.589971e-07 1.845019e-01  
## [6,] 2.715011e-07 -1.803784e-07 -1.880450e-07 -1.891093e-07 -1.714338e-07  
## [7,] 2.755309e-07 -1.849901e-07 -1.886999e-07 -1.903366e-07 -1.768260e-07  
## [8,] 2.779733e-07 -1.852205e-07 -1.966657e-07 -1.959678e-07 -1.792376e-07  
## [9,] 2.763165e-07 -1.876099e-07 -1.927455e-07 -1.909353e-07 -1.716600e-07  
## [10,] 8.986119e-07 -2.653827e-07 -2.749393e-07 -2.766843e-07 -2.509904e-07  
##           [,6]           [,7]           [,8]           [,9]  
## [1,] 5.794547e-08 5.733799e-08 5.572179e-08 5.747487e-08  
## [2,] -5.353827e-08 -5.456973e-08 -5.345037e-08 -5.575109e-08  
## [3,] -5.533367e-08 -5.442740e-08 -5.689569e-08 -5.628361e-08  
## [4,] -5.594502e-08 -5.531544e-08 -5.668502e-08 -5.554254e-08  
## [5,] -5.149345e-08 -5.296665e-08 -5.296672e-08 -5.049447e-08  
## [6,] 1.797324e-01 -5.959559e-08 -5.852159e-08 -5.810101e-08  
## [7,] -2.033618e-07 1.787948e-01 -6.033327e-08 -5.860511e-08  
## [8,] -2.039811e-07 -2.099545e-07 1.777783e-01 -5.892563e-08  
## [9,] -2.003368e-07 -2.034899e-07 -2.063762e-07 1.789836e-01  
## [10,] -2.893221e-07 -2.944453e-07 -3.010367e-07 -2.968315e-07  
##           [,10]  
## [1,] 3.238057e-07  
## [2,] -4.208164e-08
```

```
## [3,] -4.187177e-08
## [4,] -4.316316e-08
## [5,] -4.064925e-08
## [6,] -4.459324e-08
## [7,] -4.469775e-08
## [8,] -4.549216e-08
## [9,] -4.582821e-08
## [10,] 1.382470e-01
```

3


```
lambda_search = c(0, 10^(seq(-2, 2, length.out = 19))) # TODO rerun this chunk
inv_cov_matrices = aapply(lambda_search, 1, function(x) diag(get_inv_cov(x)))
colnames(inv_cov_matrices) = colnames(df)
inv_cov_matrices = cbind(lambda = lambda_search, inv_cov_matrices)

inv_cov_melted = reshape2::melt(data.frame(inv_cov_matrices),
                                   id.vars = "lambda",
                                   value.name = "inverse_covariance")

write.csv(inv_cov_melted, "inv_cov_melted.csv", row.names = FALSE)
```

```
inv_cov_melted <- read_csv("inv_cov_melted.csv")

inv_cov_melted %>%
  ggplot(aes(x = lambda, y = inverse_covariance,
             color = variable, lty = variable)) +
  geom_line() +
  scale_x_log10() +
  theme_minimal()
```



Here we can see that as lambda is increased, the values for our inverse covariance approaches zero.

4

```
kfold_cv <- function(lambda, df, folds) {
  n = nrow(df)

  # assign rows from df to different samples
  samples = sample(n, n)

  # these labels will match up to i in the for loop
  labels = rep(1:folds, each = n / folds)
  fold_labels = cbind(samples, labels)

  # pre-allocate space
  nll_vec = numeric(folds)

  for (i in 1:folds) {

    # find the rows that match i, these will be for our testing df
    hold_out <- which(fold_labels[, 2] == i)
    test_rows = fold_labels[hold_out, 1]

    # training data
```

```

train = df[-test_rows, ]

# testing (held out) data
test = df[test_rows, ]

# fit the model excluding the indices I_i
theta_hat_i = get_inv_cov(lambda, train)

# covariance computed just on the hold out
S = cov(test)

# calculate the negative log likelihood
nll_vec[i] = -log(det(theta_hat_i)) + sum(diag(S %*% theta_hat_i))
}

# computing overall negative log likelihood
return(sum(nll_vec))
}

# This also takes forever and I just saved it as a CSV so I don't have to wait
# for RMD compilation
neglik_vector = sapply(lambda_search, kfold_cv, df = df, folds = 10)
neg_lik_df <- cbind(lambda_search, neglik_vector)
write.csv(neg_lik_df, "neg_lik_df.csv", row.names = FALSE)

read_csv("neg_lik_df.csv",) %>% as.matrix()

```

```

##      lambda_search neglik_vector
## [1,]    0.00000000    -214.38992
## [2,]    0.01000000    -216.26470
## [3,]    0.01668101    -200.41873
## [4,]    0.02782559    -180.57728
## [5,]    0.04641589    -155.38860
## [6,]    0.07742637    -120.23174
## [7,]    0.12915497     -84.11231
## [8,]    0.21544347     -39.85269
## [9,]    0.35938137       2.87995
## [10,]   0.59948425      43.75154
## [11,]   1.00000000      88.93703
## [12,]   1.66810054     118.49892
## [13,]   2.78255940     150.62103
## [14,]   4.64158883     185.25581
## [15,]   7.74263683     225.00968
## [16,]  12.91549665     267.98175
## [17,]  21.54434690     314.96410
## [18,]  35.93813664     362.91954
## [19,]  59.94842503     412.14234
## [20,] 100.00000000     462.28262

```

Here we can see that the lower the value of lambda, the lower the negative log likelihood. Basically, choosing a lambda value that's 0 will result in the best results.

```

get_inv_cov_update <- function(data) {

  p = ncol(data)
  theta = Variable(rows = p, cols = p)
  S <- cov(data)
  objective = Minimize(-log_det(theta) + matrix_trace(S %*% theta))

  # Creating constraints
  i <- 2:9
  j <- 2:9

  grid <- expand.grid(i = i, j = j) %>%
    filter(i < j)

  constraint_list <- alply(grid, 1, function(row){
    # this sets everything in the upper triangle to 0
    # besides the diagonal and the upper and rightmost border of the matrix
    theta[row[, 1], row[, 2]] == 0
  })

  problem = Problem(objective, constraint_list)
  result = psolve(problem)
  return(result$getValue(theta))
}

round(get_inv_cov_update(df), 3)

##      [,1]    [,2]    [,3]    [,4]    [,5]    [,6]    [,7]    [,8]    [,9]
## [1,] 40.621    3.686  18.292    4.188  15.981   17.391   13.962    6.844   16.746
## [2,]  3.686  44.087    0.000    0.000    0.000    0.000    0.000    0.000    0.000
## [3,] 18.292    0.000  56.014    0.000    0.000    0.000    0.000    0.000    0.000
## [4,]  4.188    0.000    0.000  40.669    0.000    0.000    0.000    0.000    0.000
## [5,] 15.981    0.000    0.000    0.000  43.665    0.000    0.000    0.000    0.000
## [6,] 17.391    0.000    0.000    0.000    0.000  94.131    0.000    0.000    0.000
## [7,] 13.962    0.000    0.000    0.000    0.000    0.000  63.506    0.000    0.000
## [8,]  6.844    0.000    0.000    0.000    0.000    0.000    0.000  47.007    0.000
## [9,] 16.746    0.000    0.000    0.000    0.000    0.000    0.000    0.000  67.436
## [10,] -8.142 -16.496 -9.237 -15.319 -3.032 -30.229 -18.936 -17.909 -18.111
##      [,10]
## [1,]  -8.142
## [2,] -16.496
## [3,]  -9.237
## [4,] -15.319
## [5,]  -3.032
## [6,] -30.229
## [7,] -18.936
## [8,] -17.909
## [9,] -18.111
## [10,] 56.173

```

Here is the rounded estimate from the data.

6

```
bootstrap <- function(data) {  
  sample = sample(1:50, 50, replace = TRUE)  
  temp = data[sample, ]  
  theta_b <- get_inv_cov_update(temp)  
}  
  
# This takes forever to run, so I ran it once and saved the final result in a  
# csv down below  
B <- 100  
results <- replicate(B, bootstrap(df))
```

The bootstrap runs above. It really takes a long time to run so there won't be any output above.

```
# Constructing a grid to use to get all of the quantiles  
grid <- expand.grid(x = 1:10, y = 1:10)  
  
# this creates a dataframe for each layer of the bootstrap results  
ci_df <- adply(grid, 1, function(r){  
  quantile(results[r[,1], r[,2],], c(0.025, 0.975))  
})  
  
ci_df <- ci_df %>%  
  # filter out the values that are 0  
  filter(`2.5%` != 0 & `97.5%` != 0) %>%  
  arrange(x, y)  
  
# save the csv  
write.csv(ci_df, "ci_df.csv", row.names = FALSE)  
  
read_csv("ci_df.csv") %>% as.matrix()
```

```
##      x y      2.5%      97.5%  
## [1,] 1 1 37.053375 65.959350  
## [2,] 1 2 -6.430375 16.595400  
## [3,] 1 3  7.480150 32.502775  
## [4,] 1 4 -4.722975 17.018775  
## [5,] 1 5  9.280400 27.824325  
## [6,] 1 6  2.439400 32.165050  
## [7,] 1 7  5.586575 26.272525  
## [8,] 1 8  1.059775 19.012875  
## [9,] 1 9  6.298225 30.485350  
## [10,] 1 10 -15.715675  8.709675  
## [11,] 2 1 -6.430375 16.595400  
## [12,] 2 2 34.302400 68.435675  
## [13,] 2 10 -36.038175 -9.021975  
## [14,] 3 1  7.480150 32.502775
```

```

## [15,] 3 3 43.584575 98.534650
## [16,] 3 10 -23.896775 0.548025
## [17,] 4 1 -4.722975 17.018775
## [18,] 4 4 29.108550 59.391500
## [19,] 4 10 -25.726375 -9.178525
## [20,] 5 1 9.280400 27.824325
## [21,] 5 5 34.718450 69.420650
## [22,] 5 10 -11.711900 3.161775
## [23,] 6 1 2.439400 32.165050
## [24,] 6 6 68.942775 152.979375
## [25,] 6 10 -51.110275 -19.339050
## [26,] 7 1 5.586575 26.272525
## [27,] 7 7 49.598250 100.966325
## [28,] 7 10 -35.423775 -11.591575
## [29,] 8 1 1.059775 19.012875
## [30,] 8 8 38.827350 72.232350
## [31,] 8 10 -31.192125 -11.085925
## [32,] 9 1 6.298225 30.485350
## [33,] 9 9 50.620650 103.782775
## [34,] 9 10 -31.619475 -8.047350
## [35,] 10 1 -15.715675 8.709675
## [36,] 10 2 -36.038175 -9.021975
## [37,] 10 3 -23.896775 0.548025
## [38,] 10 4 -25.726375 -9.178525
## [39,] 10 5 -11.711900 3.161775
## [40,] 10 6 -51.110275 -19.339050
## [41,] 10 7 -35.423775 -11.591575
## [42,] 10 8 -31.192125 -11.085925
## [43,] 10 9 -31.619475 -8.047350
## [44,] 10 10 50.612475 87.058075

```

Here are all of the confidence intervals for the portion of the upper triangle that aren't zero.